

BENCHBASE FOR POSTGRES TEMPORAL QUERIES

Paul A. Jungwirth

22 August 2024

pdxpug

OUTLINE

- How to use Benchbase
- Comparing temporal foreign key implementations
- ~~More temporal procedures~~
- Benchmarking mistakes and lessons

HISTORY

OLTP-Bench: An Extensible Testbed for Benchmarking Relational Databases

Djellel Eddine Difallah
U. of Fribourg, Switzerland
djelleleddine.difallah@unifr.ch

Carlo Curino
Microsoft Corporation, USA
ccurino@microsoft.com

Andrew Pavlo
Carnegie Mellon University, USA
pavlo@cs.cmu.edu

Philippe Cudre-Mauroux
U. of Fribourg, Switzerland
pcm@unifr.ch

ABSTRACT

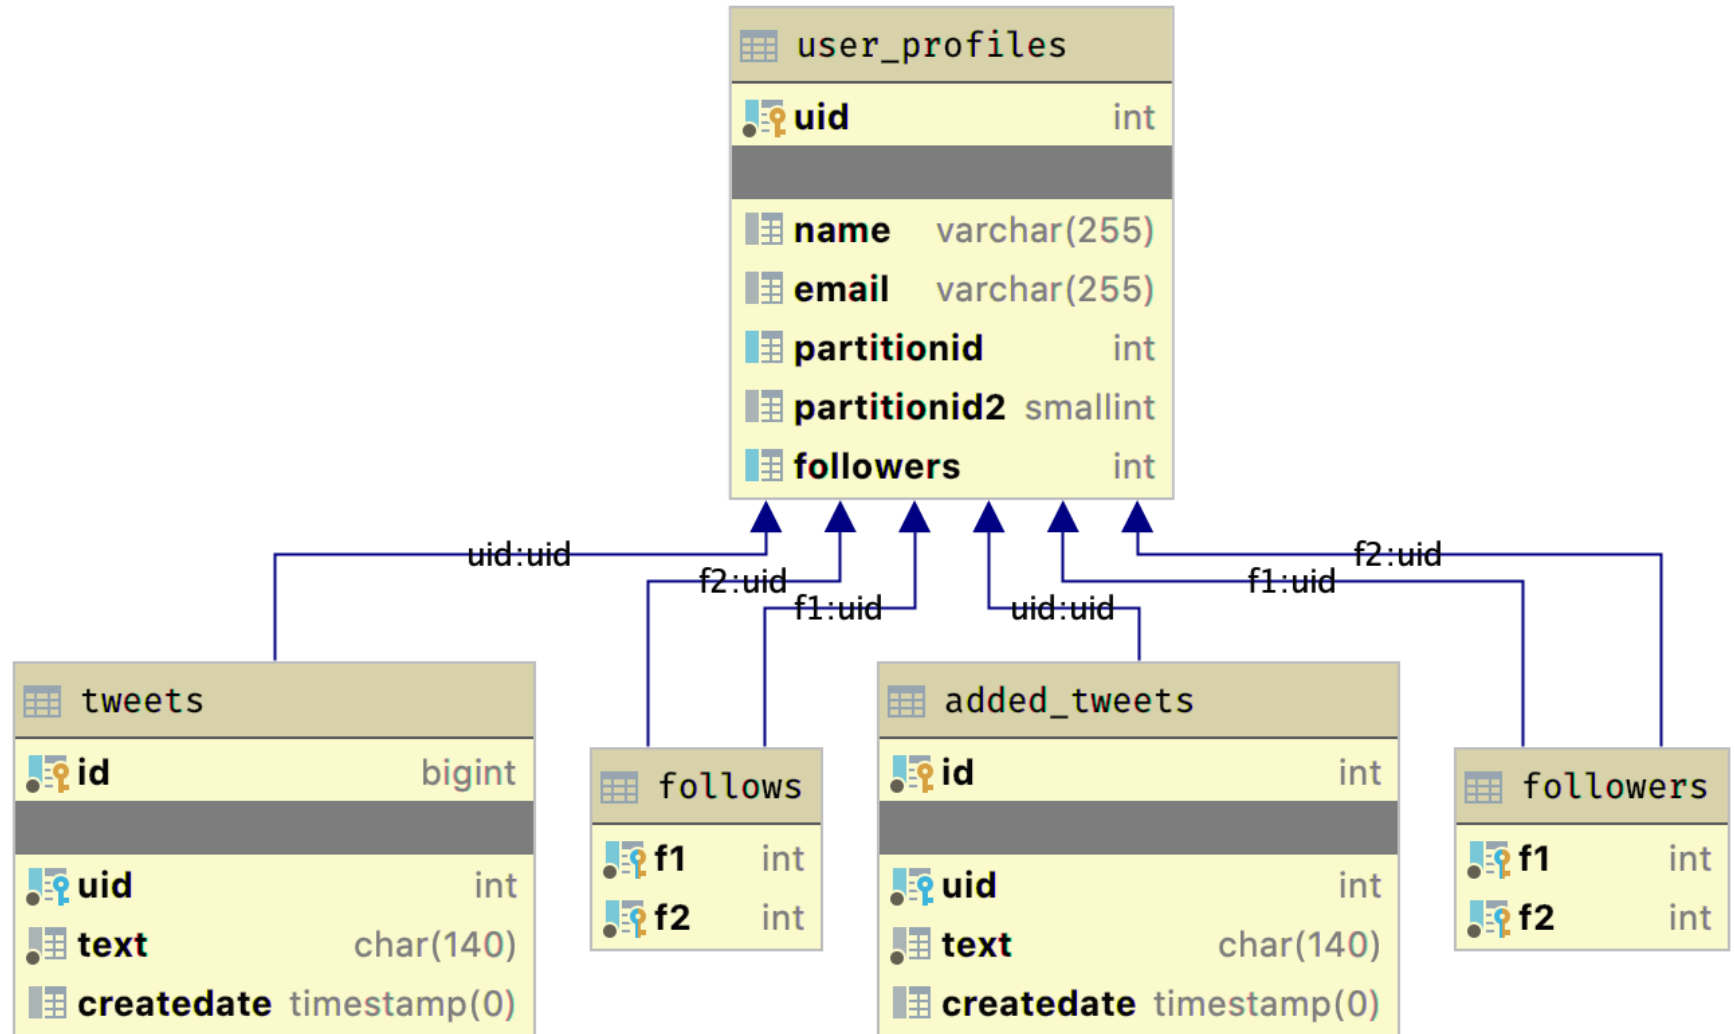
Benchmarking is an essential aspect of any database management system (DBMS) effort. Despite several recent advancements, such

To overcome this problem, it is imperative that application developers use a testing environment that is *stable*, *controlled* and *repeatable* [19]. In the context of DBMSs, this is achieved through the use of a *benchmark* that allows one to measure key performance

BENCHMARKS

```
paul@tal:~/src/benchbase$ ls \  
> src/main/java/com/oltpbenchmark/benchmarks/  
auctionmark  README.md          templated  voter  
chbenchmark  resourcestresser   temporal   wikipedia  
epinions     seats              tpcc       ycsb  
hyadapt      sibench            tpcds  
noop         smallbank          tpch  
otmetrics    tatp               twitter
```

BENCHMARKS



DBMSES

```
paul@tal:~/src/benchbase$ find \  
> src/main/resources/benchmarks/ \  
> -name 'ddl-*' | \  
> xargs -L1 basename | \  
> sort -u | \  
> pr -t -2  
ddl-cassandra.sql          ddl-nuodb.sql  
ddl-cockroachdb.sql        ddl-oracle.sql  
ddl-db2.sql                ddl-phoenix.sql  
ddl-generic.sql            ddl-postgres.sql  
ddl-hsqldb.sql             ddl-singlestore.sql  
ddl-monetdb.sql            ddl-spanner.sql  
ddl-myrocks.sql            ddl-sqlite.sql  
ddl-mysql.sql              ddl-sqlserver.sql  
ddl-noisenape.sql          ddl-timesten.sql
```

QUICKSTART

```
git clone https://github.com/cmu-db/benchbase.git
cd benchbase
./mvnw clean package -P postgres

cd target
tar xvzf benchbase-postgres.tgz
cd benchbase-postgres

java -jar benchbase.jar -b tpcc \
  -c config/postgres/sample_tpcc_config.xml \
  --create=true --load=true --execute=true
```


ONELINER

```
mvn clean compile exec:java -P postgres \  
  -Dexec.args="-b tpcc \  
  -c config/postgres/sample_tpcc_config.xml \  
  --create=true --load=true --execute=true"
```

<https://github.com/cmu-db/benchbase/pull/548>

ONELINER

```
mvn clean compile exec:java -P postgres \  
  -Dexec.args="-b tpcc \  
  -c config/postgres/sample_tpcc_config.xml \  
  --create=true --load=true --execute=true"
```

<https://github.com/cmu-db/benchbase/pull/548>

ONELINER

```
mvn clean compile exec:java -P postgres \  
  -Dexec.args="-b tpcc \  
  -c config/postgres/sample_tpcc_config.xml \  
  --create=true --load=true --execute=true"
```

<https://github.com/cmu-db/benchbase/pull/548>

ONELINER

```
mvn clean compile exec:java -P postgres \  
  -Dexec.args="-b tpcc \  
  -c config/postgres/sample_tpcc_config.xml \  
  --create=true --load=true --execute=true"
```

<https://github.com/cmu-db/benchbase/pull/548>

CONFIG FILE

```
<type>POSTGRES</type>  
<driver>org.postgresql.Driver</driver>  
<url>jdbc:postgresql://localhost:5460/benchbase?rewriteBatched  
<username>paul</username>  
<password></password>  
<reconnectOnConnectionFailure>true</reconnectOnConnectionFailure>  
<isolation>TRANSACTION_READ_COMMITTED</isolation>  
<newConnectionPerTxn>false</newConnectionPerTxn>
```

CONFIG FILE

```
<scalefactor>10000</scalefactor>  
<batchsize>128</batchsize>
```

CONFIG FILE

```
<transactiontypes>
  <transactiontype>
    <name>InsertPosition</name>
  </transactiontype>
  <transactiontype>
    <name>UpdatePosition</name>
  </transactiontype>
  <transactiontype>
    <name>UpdateEmployee</name>
  </transactiontype>
  <transactiontype>
    <name>DeleteEmployee</name>
  </transactiontype>
</transactiontypes>
```

CONFIG FILE

```
<terminals>10</terminals>  
<works>  
  <work>  
    <time>600</time>  
    <rate>10000</rate>  
    <weights>28,28,28,16</weights>  
  </work>  
</works>
```


CONFIG FILE

```
<terminals>10</terminals>
<works>
  <work arrival="poisson">
    <time>600</time>
    <rate>10000</rate>
    <weights>28,28,28,16</weights>
    <active_terminals>5</active_terminals>
  </work>
</works>
```

RESULTS

```
paul@tal:~/src/benchbase$ ls -1 results
temporal_2024-07-28_20-10-41.config.xml
temporal_2024-07-28_20-10-41.metrics.json
temporal_2024-07-28_20-10-41.params.json
temporal_2024-07-28_20-10-41.raw.csv
temporal_2024-07-28_20-10-41.results.DeleteEmployee.csv
temporal_2024-07-28_20-10-41.results.InsertPosition.csv
temporal_2024-07-28_20-10-41.results.UpdateEmployee.csv
temporal_2024-07-28_20-10-41.results.UpdatePosition.csv
temporal_2024-07-28_20-10-41.results.csv
temporal_2024-07-28_20-10-41.samples.csv
temporal_2024-07-28_20-10-41.summary.json
```

RESULTS

```
paul@tal:~/src/benchbase$ ls -1 results
temporal_2024-07-28_20-10-41.config.xml
temporal_2024-07-28_20-10-41.metrics.json
temporal_2024-07-28_20-10-41.params.json
temporal_2024-07-28_20-10-41.raw.csv
temporal_2024-07-28_20-10-41.results.DeleteEmployee.csv
temporal_2024-07-28_20-10-41.results.InsertPosition.csv
temporal_2024-07-28_20-10-41.results.UpdateEmployee.csv
temporal_2024-07-28_20-10-41.results.UpdatePosition.csv
temporal_2024-07-28_20-10-41.results.csv
temporal_2024-07-28_20-10-41.samples.csv
temporal_2024-07-28_20-10-41.summary.json
```

RESULTS

```
paul@tal:~/src/benchbase$ ls -1 results
temporal_2024-07-28_20-10-41.config.xml
temporal_2024-07-28_20-10-41.metrics.json
temporal_2024-07-28_20-10-41.params.json
temporal_2024-07-28_20-10-41.raw.csv
temporal_2024-07-28_20-10-41.results.DeleteEmployee.csv
temporal_2024-07-28_20-10-41.results.InsertPosition.csv
temporal_2024-07-28_20-10-41.results.UpdateEmployee.csv
temporal_2024-07-28_20-10-41.results.UpdatePosition.csv
temporal_2024-07-28_20-10-41.results.csv
temporal_2024-07-28_20-10-41.samples.csv
temporal_2024-07-28_20-10-41.summary.json
```

RESULTS

```
paul@tal:~/src/benchbase$ ls -1 results
temporal_2024-07-28_20-10-41.config.xml
temporal_2024-07-28_20-10-41.metrics.json
temporal_2024-07-28_20-10-41.params.json
temporal_2024-07-28_20-10-41.raw.csv
temporal_2024-07-28_20-10-41.results.DeleteEmployee.csv
temporal_2024-07-28_20-10-41.results.InsertPosition.csv
temporal_2024-07-28_20-10-41.results.UpdateEmployee.csv
temporal_2024-07-28_20-10-41.results.UpdatePosition.csv
temporal_2024-07-28_20-10-41.results.csv
temporal_2024-07-28_20-10-41.samples.csv
temporal_2024-07-28_20-10-41.summary.json
```

*.raw.csv

```
Transaction Type Index,Transaction Name,Start Time (microsecon  
4,DeleteEmployee,1722222636.414196,6786,0,1  
3,UpdateEmployee,1722222636.421013,1045,0,1  
1,InsertPosition,1722222636.422064,1402,0,1  
4,DeleteEmployee,1722222636.423471,1555,0,1  
1,InsertPosition,1722222636.425031,722,0,1  
1,InsertPosition,1722222636.425761,682,0,1
```

*.results.csv

```
Time (seconds),Throughput (requests/second),Average Latency (m
0,509.600,6.567,4.936,6.317,6.503,6.645,6.822,6.964,9.793,12.6
5,540.200,6.213,5.167,6.034,6.120,6.245,6.358,6.629,9.088,13.3
10,533.200,6.241,5.294,6.032,6.164,6.277,6.459,6.578,9.078,13.
15,539.800,6.214,2.761,5.934,6.178,6.529,6.755,7.813,12.855,30
20,739.600,4.531,2.659,3.466,3.546,5.144,6.877,7.059,11.450,58
```

*.samples.csv

```
Time (seconds),Requests,Throughput (requests/second),Minimum L
0,1780,1780.000,153,6225,6402,5613,6576,6756,7137,9928,12692
1,1794,1794.000,132,6360,6592,5582,6690,6835,6870,8501,9909
2,1795,1795.000,136,6351,6570,5568,6693,6833,6951,9773,10577
3,1803,1803.000,122,6241,6468,5531,6591,6750,6893,9495,9803
4,1901,1901.000,95,6035,6216,5249,6318,6438,6497,7804,10805
```


*.summary.json

```
{  
  "Start timestamp (milliseconds)": 1723441031104,  
  "Current Timestamp (milliseconds)": 1723441634718,  
  "Elapsed Time (nanoseconds)": 600000068387,  
  ...  
  "Measured Requests": 1147862,  
  "Latency Distribution": {  
    "95th Percentile Latency (microseconds)": 7405,  
    "Maximum Latency (microseconds)": 58595,  
    "Median Latency (microseconds)": 6198,  
    "Minimum Latency (microseconds)": 80,  
    "25th Percentile Latency (microseconds)": 5135,  
    "90th Percentile Latency (microseconds)": 6636,  
    "99th Percentile Latency (microseconds)": 11254,  
    "75th Percentile Latency (microseconds)": 6411
```

ADVANCED MONITORING

```
--monitor-type=advanced  
--interval-monitor=1000
```

ADVANCED MONITORING

```
paul@tal:~/src/benchbase$ cut -d, -f1,2,6,9 \  
> results/monitor/repeated_query_event_601.csv |  
> sed -n '1p; 2p; 3p; $p' |  
> csvtool readable -  
QueryId          Instant                               total_exec_time    execution_co  
insertPosition 2024-08-12T05:37:12.111499381Z 187.07314000000008 494  
deleteEmployee 2024-08-12T05:37:12.111499381Z 187.07314000000008 494  
updatePosition 2024-08-12T05:47:11.112707430Z 187.07314000000008 494
```

ADVANCED MONITORING

```
paul@tal:~/src/benchbase$ cut -d, -f1,2,6,9 \  
> results/monitor/repeated_query_event_601.csv |  
> sed -n '1p; 2p; 3p; $p' |  
> csvtool readable -  
QueryId          Instant          total_exec_time  execution_co  
insertPosition 2024-08-12T05:37:12.111499381Z 187.07314000000008 494  
deleteEmployee 2024-08-12T05:37:12.111499381Z 187.07314000000008 494  
updatePosition 2024-08-12T05:47:11.112707430Z 187.07314000000008 494
```

ADVANCED MONITORING

```
paul@tal:~/src/benchbase$ cut -d, -f1,2,6,9 \  
> results/monitor/repeated_query_event_601.csv |  
> sed -n '1p; 2p; 3p; $p' |  
> csvtool readable -  
QueryId          Instant          total_exec_time  execution_co  
insertPosition 2024-08-12T05:37:12.111499381Z 187.07314000000008 494  
deleteEmployee 2024-08-12T05:37:12.111499381Z 187.07314000000008 494  
updatePosition 2024-08-12T05:47:11.112707430Z 187.07314000000008 494
```

BENCHMARK ARCHITECTURE

```
paul@tal:~/src/benchbase$ ls -1 \  
> src/main/java/com/oltpbenchmark/benchmarks/temporal/  
DateRange.java  
Employee.java  
Position.java  
TemporalBenchmark.java  
TemporalConfiguration.java  
TemporalConstants.java  
TemporalLoader.java  
TemporalModel.java  
TemporalWorker.java  
procedures
```

BENCHMARK ARCHITECTURE

```
paul@tal:~/src/benchbase$ ls -1 \  
> src/main/java/com/oltpbenchmark/benchmarks/temporal/  
DateRange.java  
Employee.java  
Position.java  
TemporalBenchmark.java  
TemporalConfiguration.java  
TemporalConstants.java  
TemporalLoader.java  
TemporalModel.java  
TemporalWorker.java  
procedures
```

BENCHMARK ARCHITECTURE

```
paul@tal:~/src/benchbase$ ls -1 \  
> src/main/java/com/oltpbenchmark/benchmarks/temporal/  
DateRange.java  
Employee.java  
Position.java  
TemporalBenchmark.java  
TemporalConfiguration.java  
TemporalConstants.java  
TemporalLoader.java  
TemporalModel.java  
TemporalWorker.java  
procedures
```


BENCHMARK ARCHITECTURE

```
paul@tal:~/src/benchbase$ ls -1 \  
> src/main/java/com/oltpbenchmark/benchmarks/temporal/  
DateRange.java  
Employee.java  
Position.java  
TemporalBenchmark.java  
TemporalConfiguration.java  
TemporalConstants.java  
TemporalLoader.java  
TemporalModel.java  
TemporalWorker.java  
procedures
```

BENCHMARK ARCHITECTURE

```
paul@tal:~/src/benchbase$ ls -1 \  
> src/main/java/com/oltpbenchmark/benchmarks/temporal/  
DateRange.java  
Employee.java  
Position.java  
TemporalBenchmark.java  
TemporalConfiguration.java  
TemporalConstants.java  
TemporalLoader.java  
TemporalModel.java  
TemporalWorker.java  
procedures
```

BENCHMARK ARCHITECTURE

```
paul@tal:~/src/benchbase$ ls -1 \  
> src/main/java/com/oltpbenchmark/benchmarks/temporal/  
DateRange.java  
Employee.java  
Position.java  
TemporalBenchmark.java  
TemporalConfiguration.java  
TemporalConstants.java  
TemporalLoader.java  
TemporalModel.java  
TemporalWorker.java  
procedures
```

BENCHMARK ARCHITECTURE

```
paul@tal:~/src/benchbase$ ls -1 \  
> src/main/java/com/oltpbenchmark/benchmarks/temporal/  
DateRange.java  
Employee.java  
Position.java  
TemporalBenchmark.java  
TemporalConfiguration.java  
TemporalConstants.java  
TemporalLoader.java  
TemporalModel.java  
TemporalWorker.java  
procedures
```

TemporalWorker

```
@Override
protected TransactionStatus executeWork(
    Connection conn, TransactionType nextTrans)
    throws UserAbortException, SQLException {
    try {
        if (nextTrans.getProcedureClass().equals(InsertPosition.class))
            RandomEmployee emp = makeRandomEmployee(
                TemporalConstants.CHECK_FK_GAUSSIAN_RANGE,
                config.getMaxYearsInsertPositionRange());
        String duty = TemporalConstants.POSITION_NAMES[
            rng().nextInt(TemporalConstants.POSITION_NAMES.length)];
        int rank = 1;

        getProcedure(InsertPosition.class).run(conn, emp.id, duty,
```

TemporalWorker

```
@Override
protected TransactionStatus executeWork(
    Connection conn, TransactionType nextTrans)
    throws UserAbortException, SQLException {
    try {
        if (nextTrans.getProcedureClass().equals(InsertPosition.class))
            RandomEmployee emp = makeRandomEmployee(
                TemporalConstants.CHECK_FK_GAUSSIAN_RANGE,
                config.getMaxYearsInsertPositionRange());
        String duty = TemporalConstants.POSITION_NAMES[
            rng().nextInt(TemporalConstants.POSITION_NAMES.length)];
        int rank = 1;

        getProcedure(InsertPosition.class).run(conn, emp.id, duty,
```

TemporalWorker

```
@Override
protected TransactionStatus executeWork(
    Connection conn, TransactionType nextTrans)
    throws UserAbortException, SQLException {
    try {
        if (nextTrans.getProcedureClass().equals(InsertPosition.class)) {
            RandomEmployee emp = makeRandomEmployee(
                TemporalConstants.CHECK_FK_GAUSSIAN_RANGE,
                config.getMaxYearsInsertPositionRange());
            String duty = TemporalConstants.POSITION_NAMES[
                rng().nextInt(TemporalConstants.POSITION_NAMES.length)];
            int rank = 1;

            getProcedure(InsertPosition.class).run(conn, emp.id, duty,
```

TemporalWorker

```
connection conn; TransactionType nextTrans;
throws UserAbortException, SQLException {
try {
    if (nextTrans.getProcedureClass().equals(InsertPosition.class))
        RandomEmployee emp = makeRandomEmployee(
            TemporalConstants.CHECK_FK_GAUSSIAN_RANGE,
            config.getMaxYearsInsertPositionRange());
        String duty = TemporalConstants.POSITION_NAMES[
            rng().nextInt(TemporalConstants.POSITION_NAMES.length)];
        int rank = 1;

        getProcedure(InsertPosition.class).run(conn, emp.id, duty, rank);

    } else if (nextTrans.getProcedureClass().equals(UpdatePosition.class))
        RandomPosition p = makeRandomPosition(config.getMaxYearsUpdatePositionRange());
}
```


TemporalWorker

```
@Override
protected TransactionStatus executeWork(
    Connection conn, TransactionType nextTrans)
    throws UserAbortException, SQLException {
    try {
        if (nextTrans.getProcedureClass().equals(InsertPosition.class))
            RandomEmployee emp = makeRandomEmployee(
                TemporalConstants.CHECK_FK_GAUSSIAN_RANGE,
                config.getMaxYearsInsertPositionRange());
            String duty = TemporalConstants.POSITION_NAMES[
                rng().nextInt(TemporalConstants.POSITION_NAMES.length)];
            int rank = 1;

            getProcedure(InsertPosition.class).run(conn, emp.id, duty,
                rank);
        } else if (nextTrans.getProcedureClass().equals(UpdatePosition.class)) {
            RandomEmployee emp = makeRandomEmployee(
                TemporalConstants.CHECK_FK_GAUSSIAN_RANGE,
                config.getMaxYearsInsertPositionRange());
            String duty = TemporalConstants.POSITION_NAMES[
                rng().nextInt(TemporalConstants.POSITION_NAMES.length)];
            int rank = 1;

            getProcedure(UpdatePosition.class).run(conn, emp.id, duty,
                rank);
        }
    } catch (SQLException e) {
        throw new UserAbortException(e);
    }
}
```

PROCEDURES

```
public class InsertPosition extends Procedure {
    public final SQLStmt insertPosition =
        new SQLStmt(
            "INSERT INTO positions (employee_id, valid_at, name) " +
            "VALUES (?, daterange(?, ?), " +
            "concat(?, ' ', to_char(?, 'RN')) RETURNING id");

    public int run(
        Connection conn, int employeeId, String duty,
        LocalDate s, LocalDate e, int rank)
        throws SQLException {
        try (PreparedStatement stmt = this.getPreparedStatement(conn,
            stmt.setInt(1, employeeId);
            stmt.setDate(2, s == null ? null : Date.valueOf(s));
            stmt.setDate(3, e == null ? null : Date.valueOf(e));
```

PROCEDURES

```
public class InsertPosition extends Procedure {
    public final SQLStmt insertPosition =
        new SQLStmt(
            "INSERT INTO positions (employee_id, valid_at, name) " +
            "VALUES (?, daterange(?, ?), " +
            "concat(?, ' ', to_char(?, 'RN'))) RETURNING id");

    public int run(
        Connection conn, int employeeId, String duty,
        LocalDate s, LocalDate e, int rank)
        throws SQLException {
        try (PreparedStatement stmt = this.getPreparedStatement(conn,
            stmt.setInt(1, employeeId);
            stmt.setDate(2, s == null ? null : Date.valueOf(s));
            stmt.setDate(3, e == null ? null : Date.valueOf(e));
```

PROCEDURES

```
public class InsertPosition extends Procedure {
    public final SQLStmt insertPosition =
        new SQLStmt(
            "INSERT INTO positions (employee_id, valid_at, name) " +
            "VALUES (?, daterange(?, ?), " +
            "concat(?, ' ', to_char(?, 'RN'))) RETURNING id");

    public int run(
        Connection conn, int employeeId, String duty,
        LocalDate s, LocalDate e, int rank)
        throws SQLException {
        try (PreparedStatement stmt = this.getPreparedStatement(conn,
            insertPosition)) {
            stmt.setInt(1, employeeId);
            stmt.setDate(2, s == null ? null : Date.valueOf(s));
            stmt.setDate(3, e == null ? null : Date.valueOf(e));
        }
    }
}
```

TESTS

```
paul@tal:~/src/benchbase$ ls \
> src/test/java/com/oltpbenchmark/benchmarks/
auctionmark  resourcestresser  temporal  wikipedia
chbenchmark  seats             tpcc      ycsb
epinions     smallbank        tpch
noop         tatp             twitter
otmetrics    templated        voter
```

TEMPORAL DDL

```
DROP EXTENSION IF EXISTS btree_gist CASCADE;
DROP TABLE IF EXISTS employees CASCADE;
DROP TABLE IF EXISTS positions CASCADE;

CREATE EXTENSION btree_gist;

CREATE TABLE employees (
  id          int GENERATED BY DEFAULT AS IDENTITY NOT NULL,
  valid_at    daterange NOT NULL,
  name        text NOT NULL,
  salary      int NOT NULL,
  PRIMARY KEY (id, valid_at WITHOUT OVERLAPS)
);

CREATE TABLE positions (
```

TEMPORAL DDL

```
DROP TABLE IF EXISTS positions CASCADE;
```

```
CREATE EXTENSION btree_gist;
```

```
CREATE TABLE employees (  
  id          int GENERATED BY DEFAULT AS IDENTITY NOT NULL,  
  valid_at    daterange NOT NULL,  
  name        text NOT NULL,  
  salary      int NOT NULL,  
  PRIMARY KEY (id, valid_at WITHOUT OVERLAPS)  
);
```

```
CREATE TABLE positions (  
  id          int GENERATED BY DEFAULT AS IDENTITY NOT NULL,  
  valid at    daterange NOT NULL,
```

TEMPORAL DDL

```
salary      int NOT NULL,  
PRIMARY KEY (id, valid_at WITHOUT OVERLAPS)  
);  
  
CREATE TABLE positions (  
  id          int GENERATED BY DEFAULT AS IDENTITY NOT NULL,  
  valid_at    daterange NOT NULL,  
  name        text NOT NULL,  
  employee_id int NOT NULL,  
  PRIMARY KEY (id, valid_at WITHOUT OVERLAPS),  
  FOREIGN KEY (employee_id, PERIOD valid_at)  
    REFERENCES employees (id, PERIOD valid_at)  
);  
CREATE INDEX idx_positions_employee_id ON positions  
  USING gist (employee_id, valid_at);
```



TEMPORAL DDL

```
salary      int NOT NULL,  
PRIMARY KEY (id, valid_at WITHOUT OVERLAPS)  
);  
  
CREATE TABLE positions (  
  id          int GENERATED BY DEFAULT AS IDENTITY NOT NULL,  
  valid_at    daterange NOT NULL,  
  name        text NOT NULL,  
  employee_id int NOT NULL,  
  PRIMARY KEY (id, valid_at WITHOUT OVERLAPS),  
  FOREIGN KEY (employee_id, PERIOD valid_at)  
    REFERENCES employees (id, PERIOD valid_at)  
);  
CREATE INDEX idx_positions_employee_id ON positions  
  USING gist (employee_id, valid_at);
```

TEMPORAL DDL

```
salary      int NOT NULL,  
PRIMARY KEY (id, valid_at WITHOUT OVERLAPS)  
);  
  
CREATE TABLE positions (  
  id          int GENERATED BY DEFAULT AS IDENTITY NOT NULL,  
  valid_at    daterange NOT NULL,  
  name        text NOT NULL,  
  employee_id int NOT NULL,  
  PRIMARY KEY (id, valid_at WITHOUT OVERLAPS),  
  FOREIGN KEY (employee_id, PERIOD valid_at)  
    REFERENCES employees (id, PERIOD valid_at)  
);  
  
CREATE INDEX idx_positions_employee_id ON positions  
  USING gist (employee_id, valid_at);
```

TEMPORAL FOREIGN KEYS



range_agg

IMPLEMENTATION

```
SELECT 1
FROM (
  SELECT pkperiodatt AS r
  FROM [ONLY] pktable x
  WHERE pkatt1 = $1 [AND ...]
  AND pkperiodatt && $n
  FOR KEY SHARE OF x
) x1
HAVING $n <@ range_agg(x1.r)
```

EXISTS

IMPLEMENTATION

```
SELECT 1
-- There was a PK when the FK started:
WHERE EXISTS
    SELECT 1
    FROM [ONLY] <pktable>
    WHERE pkatt1 = $1 [AND ...]
    AND COALESCE(lower(pkperiodatt), '-Infinity')
        <= COALESCE(lower($n), '-Infinity')
    AND COALESCE(lower($n), '-Infinity')
        < COALESCE(upper(pkperiodatt), 'Infinity')
)
...
```

EXISTS

IMPLEMENTATION

```
-- There was a PK when the FK ended:  
AND EXISTS (  
  SELECT 1  
  FROM    [ONLY] <pktable>  
  WHERE    pkatt1 = $1 [AND ...]  
  AND      COALESCE(lower(pkperiodatt), '-Infinity')  
            < COALESCE(upper($n), 'Infinity')  
  AND      COALESCE(upper($n), 'Infinity')  
            <= COALESCE(upper(pkperiodatt), 'Infinity')  
)  
...
```

EXISTS

IMPLEMENTATION

```
-- There are no gaps in the PK:  
-- (i.e. there is no PK that ends early,  
-- unless a matching PK record starts right away)  
AND NOT EXISTS (  
  SELECT 1  
  FROM    [ONLY] <pktable> AS pk1  
  WHERE   pkatt1 = $1 [AND ...]  
  AND     COALESCE(lower($n), '-Infinity')  
          < COALESCE(upper(pkperiodatt), 'Infinity')  
  AND     COALESCE(upper(pkperiodatt), 'Infinity')  
          < COALESCE(upper($n), 'Infinity')  
  ...  
)
```

EXISTS

IMPLEMENTATION

```
AND      NOT EXISTS (  
  SELECT 1  
  FROM    [ONLY] <pktable> AS pk2  
  WHERE   pk1.pkatt1 = pk2.pkatt1 [AND ...]  
          -- but skip pk1.pkperiodatt && pk2.pkperiodatt  
  AND     COALESCE(lower(pk2.pkperiodatt), '-Infinity')  
          <= COALESCE(upper(pk1.pkperiodatt), 'Infinity')  
          COALESCE(upper(pk1.pkperiodatt), 'Infinity')  
          <  COALESCE(upper(pk2.pkperiodatt), 'Infinity')  
  )  
)
```


Lag

IMPLEMENTATION

```
SELECT 1
FROM (
  SELECT uk.uk_start_value,
         uk.uk_end_value,
         NULLIF(LAG(uk.uk_end_value) OVER (ORDER BY uk.uk_sta
FROM (
  SELECT coalesce(lower(x.pkperiodatt), '-Infinity') AS uk_
         coalesce(upper(x.pkperiodatt), 'Infinity') AS uk_e
  FROM   pktable AS x
  WHERE  pkatt1 = $1 [AND ...]
  AND    uk.pkperiodatt && $n
  FOR KEY SHARE OF x
) AS uk
) AS uk
WHERE uk.uk_start_value < inner($n)
```

COMPILER FLAGS

```
#if defined(RI_TEMPORAL_IMPL_LAG)
    quoteOneName(attname,

    appendStringInfo(&querybuf, "SELECT 1 FROM ( ");
    appendStringInfo(&querybuf, "    SELECT    uk.uk_start_value,
    appendStringInfo(&querybuf, "                NULLIF(LAG(uk.uk_en
    appendStringInfo(&querybuf, "    FROM      ( ");
    appendStringInfo(&querybuf, "        SELECT    COALESCE(LOWER(x.
    appendStringInfo(&querybuf, "                COALESCE(UPPER(x.
    appendStringInfo(&querybuf, "        FROM      %s%s AS x", pk_on

#elif defined(RI_TEMPORAL_IMPL_EXISTS)
    appendStringInfo(&querybuf,
                    "SELECT 1 ");

#else
    quoteOneName(attname
```

EXPLAIN range_agg

```
Aggregate
  Filter: ('[2020-10-10,2020-12-12)')::daterange <@ range_agg(x
    -> Subquery Scan on x1
      -> LockRows
        -> Index Scan using employees_pkey on employees x
          Index Cond: ((id = 500) AND (valid_at && '[2020-10-10,
```

EXPLAIN Lag

Aggregate

Filter: ((array_agg(uk.x) FILTER (WHERE (uk.x IS NOT NULL)))

-> Subquery Scan on uk

Filter: ((uk.uk_start_value < '2020-12-12'::date) AND (uk.

-> WindowAgg

-> Sort

Sort Key: uk_1.uk_start_value

-> Subquery Scan on uk_1

-> LockRows

-> Index Scan using employees_pkey on employees x

Index Cond: ((id = 500) AND (valid_at && '[2020-

EXPLAIN EXISTS

Result

One-Time Filter: ((InitPlan 1).col1 AND (InitPlan 2).col1 AND

InitPlan 1

-> LockRows

-> Index Scan using employees_pkey on employees x

Index Cond: ((id = 500) AND (valid_at && '[2020-10-10,20

Filter: ((COALESCE(lower(valid_at), '-infinity'::date) <

InitPlan 2

-> LockRows

-> Index Scan using employees_pkey on employees x_1

Index Cond: ((id = 500) AND (valid_at && '[2020-10-10,20

Filter: ((COALESCE(lower(valid_at), '-infinity'::date) <

InitPlan 4

-> LockRows

-> Index Scan using employees_pkey on employees nk1

ASSUMPTIONS

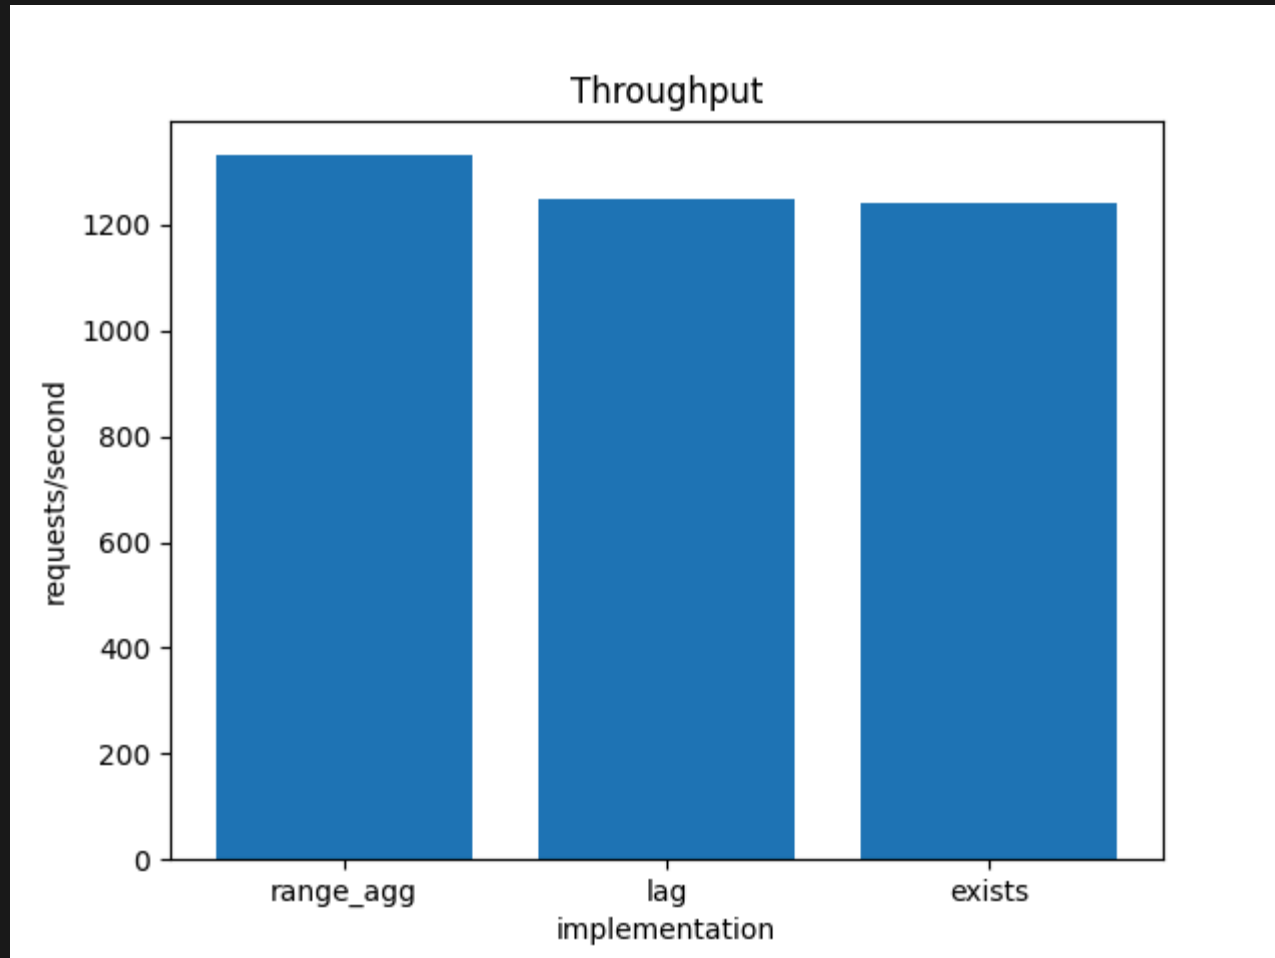
Relevant:

- CPU
- Tuples examined
- Number of index scans

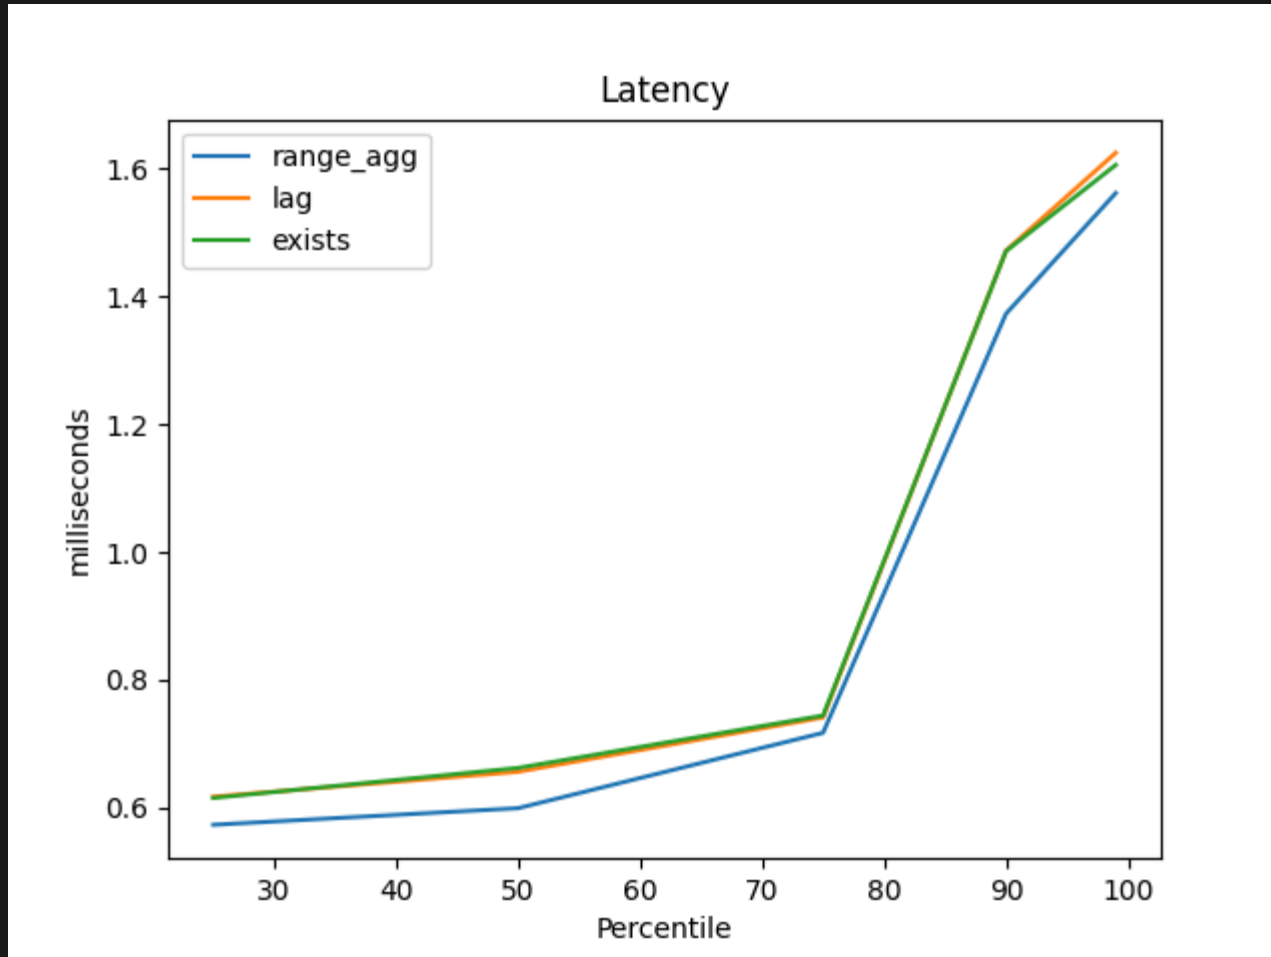
Not relevant:

- Shared Buffers
- I/O

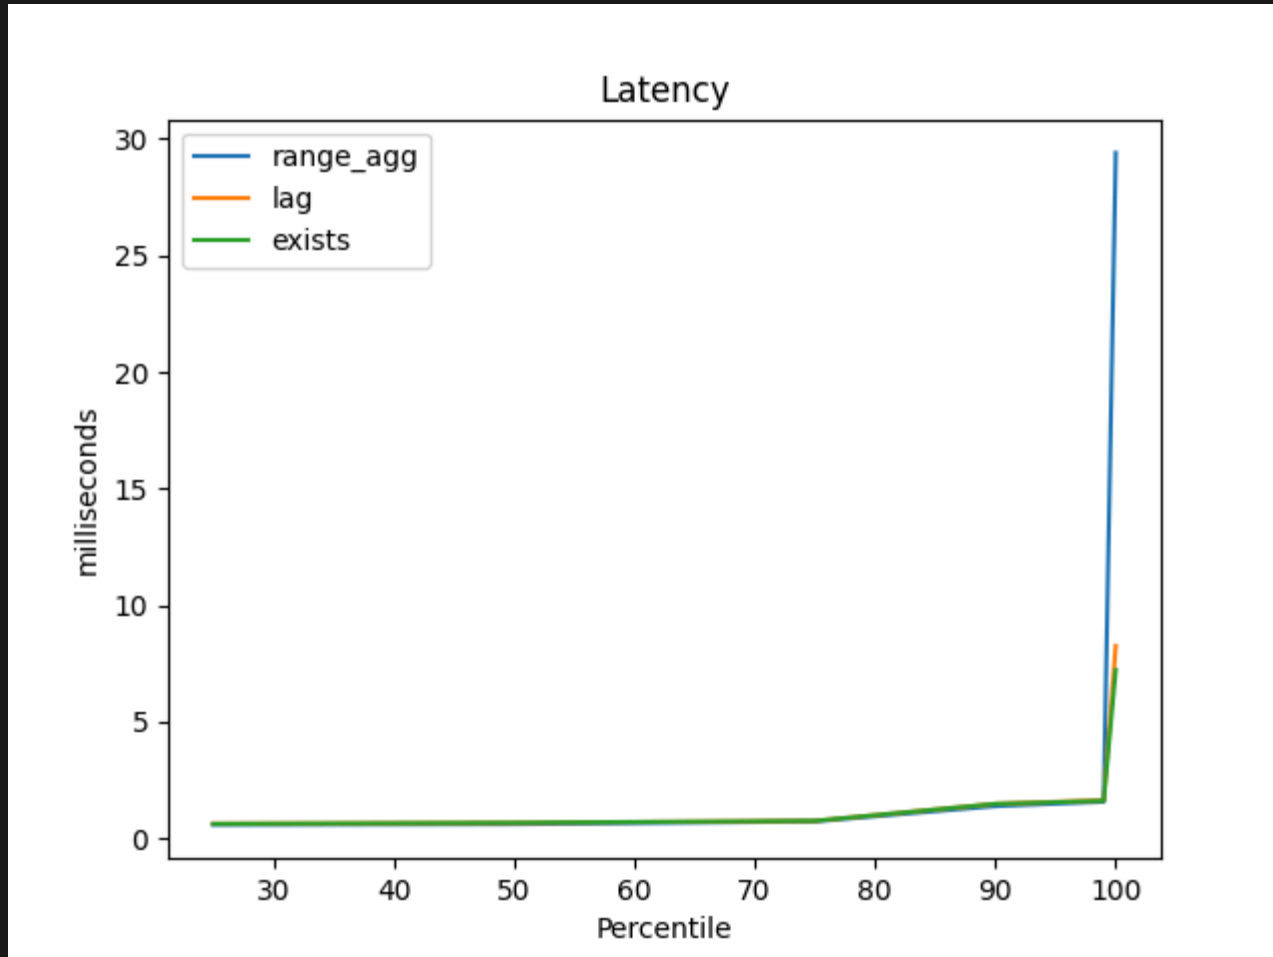
EARLY RESULTS



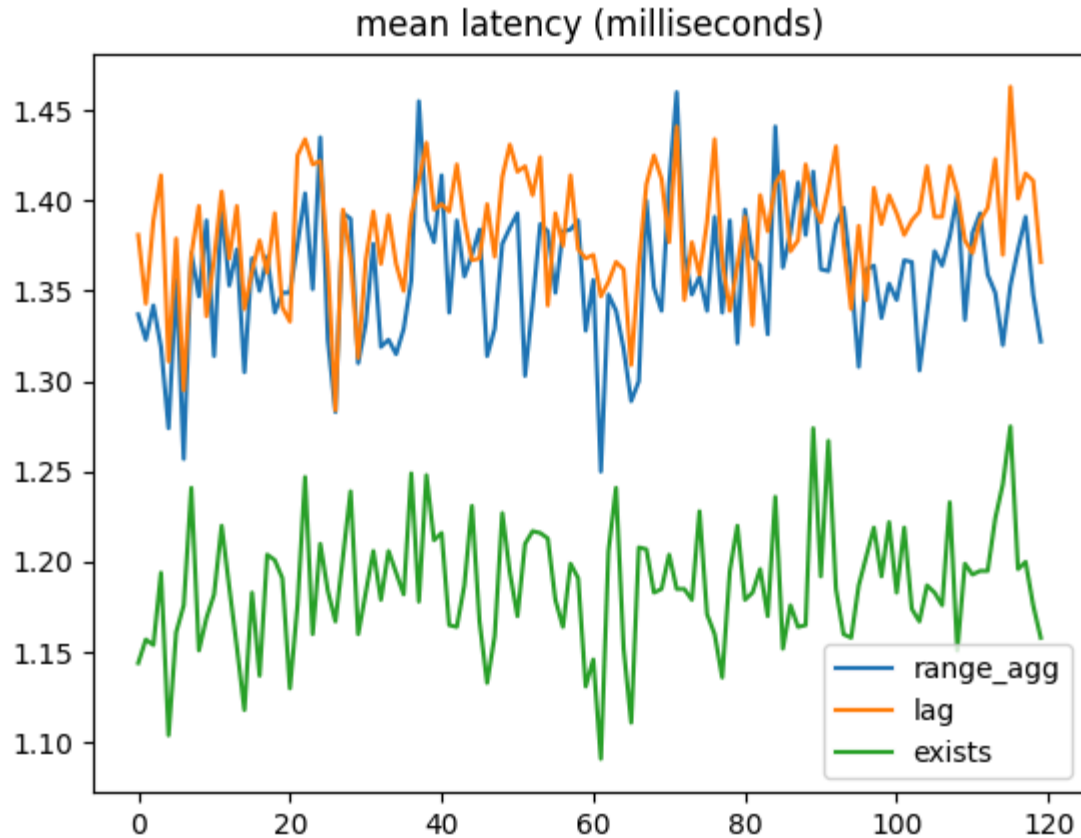
EARLY RESULTS




EARLY RESULTS



SURPRISE



SURPRISE



50% ERRORS

Completed Transactions:

com.oltpbenchmark.benchmarks.temporal.procedures.CheckForeignKeyRangeAgg/01	[72064]	*****
com.oltpbenchmark.benchmarks.temporal.procedures.CheckForeignKeyLag/02	[71479]	*****
com.oltpbenchmark.benchmarks.temporal.procedures.CheckForeignKeyExists/03	[71529]	*****
com.oltpbenchmark.benchmarks.temporal.procedures.Noop/04	[4585]	*****

Aborted Transactions:

<EMPTY>

Rejected Transactions (Server Retry):

<EMPTY>

Rejected Transactions (Retry Different):

<EMPTY>

Unexpected SQL Errors:

com.oltpbenchmark.benchmarks.temporal.procedures.CheckForeignKeyRangeAgg/01	[80861]	*****
com.oltpbenchmark.benchmarks.temporal.procedures.CheckForeignKeyLag/02	[80764]	*****
com.oltpbenchmark.benchmarks.temporal.procedures.CheckForeignKeyExists/03	[80478]	*****

EXPLAIN ANALYZE EXISTS

```
Result (actual time=0.034..0.035 rows=0 loops=1)
  One-Time Filter: ((InitPlan 1).col1 AND (InitPlan 2).col1 AND (InitPlan 3).col1)
  InitPlan 1
    -> LockRows (actual time=0.033..0.033 rows=0 loops=1)
      -> Index Scan using employees_pkey on employees x (actual time=0.033..0.033 rows=0 loops=1)
        Index Cond: ((id = 5999) AND (valid_at && '[2020-10-10,2020-10-11)' AND (COALESCE(lower(valid_at), '-infinity'::date) < '2020-10-10'))
        Filter: ((COALESCE(lower(valid_at), '-infinity'::date) < '2020-10-10'))
      InitPlan 2
        -> LockRows (never executed)
          -> Index Scan using employees_pkey on employees x_1 (never executed)
            Index Cond: ((id = 5999) AND (valid_at && '[2020-10-10,2020-10-11)' AND (COALESCE(lower(valid_at), '-infinity'::date) < '2020-10-10'))
            Filter: ((COALESCE(lower(valid_at), '-infinity'::date) < '2020-10-10'))
          InitPlan 4
            -> LockRows (never executed)
              -> Index Scan using employees_pkey on employees nk1 (never executed)
```

EXPLAIN ANALYZE EXISTS

```
Filter: ((COALESCE(lower(valid_at), '-infinity'::date) <
InitPlan 2
-> LockRows (never executed)
-> Index Scan using employees_pkey on employees x_1 (never
Index Cond: ((id = 5999) AND (valid_at && '[2020-10-10,2
Filter: ((COALESCE(lower(valid_at), '-infinity'::date) <
InitPlan 4
-> LockRows (never executed)
-> Index Scan using employees_pkey on employees pk1 (never
Index Cond: ((id = 5999) AND (valid_at && '[2020-10-10,2
Filter: (('2020-10-10'::date < COALESCE(upper(valid_at),
SubPlan 3
-> LockRows (never executed)
-> Index Scan using employees_pkey on employees pk2 (
Index Cond: (id = pk1.id)
```

bpfttrace ExecProcNode

```
// Count how many exec nodes per query were required,  
// and print a histogram of how often each count happens.  
// Run this for each FK implementation separately.  
// My hypothesis is that the EXISTS implementation calls ExecP  
// but only if the FK is invalid.
```

```
u:/home/paul/local/bench-*/bin/postgres:standard_ExecutorStart  
    @nodes[tid] = 0  
}  
u:/home/paul/local/bench-*/bin/postgres:ExecProcNode {  
    @nodes[tid] += 1  
}  
u:/home/paul/local/bench-*/bin/postgres:standard_ExecutorEnd {  
    @calls = hist(@nodes[tid]);  
    delete(@nodes[tid]).
```

bpfttrace ExecProcNode

all valid:

```
@calls:
```

[0]	6
-----	---

[1]	0
-----	---

$[2, 4)$	\emptyset
----------	-------------

```
[4, 8)      228851 | @@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@
```

$[8, 16)$	1
-----------	---

[16, 32)	1
----------	---

[32, 64)	2
----------	---

[64, 128)	2
-----------	---

[128, 256) 2

[256, 512)	5
------------	---


bpfttrace ExecProcNode

50+% invalid:


@calls:

[0]	6	
[1]	0	
[2, 4)	218294	@@@@
[4, 8)	183438	@@@@
[8, 16)	231	
[16, 32)	1	
[32, 64)	2	
[64, 128)	2	
[128, 256)	2	
[256, 512)	5	


MOSTLY VALID



MOSTLY VALID



MOSTLY VALID



METHODOLOGY

- Short iterations
- Automate
- Keep notes
- Automating keeping notes?

NEXT

- More questions
- More procedures
- How to distribute the benchmark?

THANKS!

REFERENCES

THIS TALK

- This talk: <https://github.com/pjungwir/pdpxpug2024-benchbase-and-temporal-foreign-keys>
- Temporal benchmark on my Basebase branch: <https://github.com/pjungwir/benchbase/tree/temporal>
- Benchmark notes and tools: <https://github.com/pjungwir/benchmarking-temporal-tables>

BENCHBASE

- Djellel Eddine Difallah and Andrew Pavlo and Carlo Curino and Philippe Cudré-Mauroux, "OLTP-Bench: An Extensible Testbed for Benchmarking Relational Databases," *PVLDB* 7.4, 2013, <http://www.vldb.org/pvldb/vol7/p277-difallah.pdf>
- Original OLTP-Bench repo wiki: <https://github.com/oltpbenchmark/oltpbench/wiki>
- Benchbase repo: <https://github.com/cmu-db/benchbase>
- PR to fix the exec:java target: <https://github.com/cmu-db/benchbase/pull/548>

TEMPORAL

- Richard Snodgrass, *Developing Time-Oriented Applications in SQL*, <https://www2.cs.arizona.edu/~rts/tdbbook.pdf>
- periods extension: <https://github.com/xocolatl/periods>
- My temporal patches: <https://commitfest.postgresql.org/49/4308/>
- My temporal branch: <https://github.com/pjungwir/postgresql/tree/valid-time>
- My temporal FK comparison branch: <https://github.com/pjungwir/postgresql/tree/temporal-fk-comparison>

BENCHMARKING

- Andres Freund, "Analyzing Postgres performance problems using perf and eBPF," <https://www.youtube.com/watch?v=HghP4D72Noc>
- Claire Giordano, "How I got started as a developer (& in Postgres), with Melanie Plageman & Thomas Munro," PathToCitusCon episode 4, <https://www.youtube.com/watch?v=72OdrpZXjEg>
- Mark Callaghan, Small Datum, <https://smalldatum.blogspot.com>
- Melanie Plageman, "Visualizing Postgres I/O Performance," PGCon 2023, <https://www.youtube.com/watch?v=CxyPZHG5bel>
- Melanie Plageman, "Postgres Performance Observability Sources and Analysis Techniques," <https://www.youtube.com/watch?v=laxZdbE1Nuw>
- Michael Christofides and Nikolay Samokhvalov, "Getting started with benchmarking," Postgres.FM episode 110, <https://www.youtube.com/watch?v=xR-VJjR9DPQ>